

SAIKIRAN S

Azure Data Engineer | ETL/ELT Specialist | Azure Data Factory | Databricks | Azure Synapse | Delta Lake

📞 +91 9591495564 | ✉️ saikiransofficial@gmail.com | 📍 Bengaluru, India

PROFESSIONAL SUMMARY

Results driven Azure Data Engineer with 4+ years of experience designing and delivering enterprise-scale data platforms across pharmaceutical, healthcare, and retail domains. Experienced in building scalable, reliable, and cost-efficient data pipelines using Azure Data Factory, Azure Databricks, PySpark, Azure Synapse Analytics, Delta Lake, ADLS Gen2, and Azure SQL Database. Strong expertise in Lakehouse architecture, Medallion (Bronze-Silver-Gold) data modeling, Common Data Models (CDM), Spark performance optimization, incremental data processing, and end-to-end data warehousing. Proven ability to optimize enterprise data platforms, streamline country onboarding, and deliver reliable, production-grade analytics solutions. Passionate about continuous learning and expanding expertise in Artificial Intelligence, Machine Learning, and modern data engineering technologies.

TECHNICAL SKILLS

Languages: Advanced SQL, Python, PySpark, Spark SQL, T-SQL

Orchestration: Azure Data Factory (ADF), ADF Triggers, Linked Services, Pipelines, Integration Runtime

Compute & Processing: Azure Databricks, Apache Spark, PySpark, Delta Live Tables, DBT (Transformations)

Storage: ADLS Gen2, Azure Blob Storage, Azure SQL Database, Amazon S3

Data Warehousing: Azure Synapse Analytics, Microsoft Fabric, Delta Lake

Big Data & Streaming: Apache Spark, PySpark, Delta Lake, Delta Live Tables (DLT), Spark Structured Streaming, Apache Kafka, Auto Loader, Unity Catalog

Architecture: Medallion (Bronze/Silver/Gold), Star Schema, Snowflake Schema, SCD Type 2, Incremental Loading

Python Libraries: Pandas, NumPy, PySpark DataFrames, Spark MLlib

Monitoring & Ops: Azure Monitor, ADF Pipeline Monitoring, Log Analytics, Azure Alerts

DevOps & VCS: Azure DevOps, Git, CI/CD Pipelines, JIRA

Security: Azure Key Vault, RBAC, Unity Catalog Governance

Testing: ETL Testing, DWH Testing, Data Validation, Reconciliation, KPI Validation

Reporting: Power BI, Azure Synapse Serverless SQL, Microsoft Excel

Data Integration: REST APIs

Streaming: Spark Structured Streaming, Apache Kafka

Methodologies: Agile, Scrum, SDLC

PROFESSIONAL EXPERIENCE

Azure Data Engineer | Solutionec | Bengaluru, India

April 2024 – April 2026

- Designed, developed, and supported scalable Azure Data Engineering pipelines for both batch and near real-time data processing using Azure Data Factory, Azure Databricks, Azure Event Hubs, Databricks Structured Streaming, Auto Loader, and Delta Lake.
- Developed reusable PySpark notebooks, Python scripts, and advanced SQL solutions implementing business transformations, data quality validation, schema standardization, incremental data processing, and enterprise data warehouse development.
- Applied Spark architecture principles, PySpark performance optimization, Delta Lake, Medallion Architecture (Bronze, Silver, Gold), Azure Synapse Analytics, Unity Catalog, and end-to-end data warehouse testing to build scalable, governed, and high-performance analytical platforms.
- Collaborated with cross-functional teams throughout the software development lifecycle to deliver enterprise Azure data platforms, contributing to a 67% reduction in reporting latency, 40% faster country onboarding, and improved production pipeline performance.

Data Analyst | Atria | Bengaluru, India

February 2022 – March 2024

- Worked primarily on enterprise Master Data Management (MDM) initiatives for multiple US pharmaceutical clients, building a strong foundation in Healthcare and Life Sciences commercial data domains including HCP, HCO, Product, Customer, and Address master data.
- Developed PySpark notebooks, Python scripts, and reusable transformation frameworks implementing business rules, data cleansing, standardization, referential integrity checks, data profiling, and enterprise MDM ingestion processes.
- Authored advanced SQL queries, database views, joins, window functions, and reconciliation frameworks for development, KPI validation, source-to-target validation, and large-scale healthcare data analysis.

- Gained hands-on experience across Azure and AWS cloud ecosystems using Azure Data Factory, Azure Databricks, Azure SQL Database, ADLS Gen2, Amazon S3, REST APIs, Azure DevOps, and JIRA while supporting enterprise data engineering and MDM integration workflows.

KEY PROJECTS

Project 1: SINERGI 2 – Market Analytics

Client: Leading US Pharmaceutical Client | Solutionec

Tech Stack: Azure Data Factory (ADF), Azure Databricks, Azure Synapse Analytics, PySpark, Spark SQL, DBT, Delta Lake, ADLS Gen2, Azure SQL Database, Unity Catalog, Azure Key Vault, Azure DevOps, Git, Power BI

- Designed and implemented a metadata-driven Azure Lakehouse solution to ingest, standardize, and harmonize multi-country pharmaceutical market and sales data into a centralized Common Data Model (CDM), enabling consistent reporting and analytics across global affiliates.
- Developed reusable Azure Data Factory pipelines and PySpark notebooks for watermark-based incremental loading, schema harmonization, recoding local business attributes into global CDM attributes, business rule implementation, market definitions, product groupings, data quality validation, deduplication, Delta MERGE/UPSERT operations, and SCD Type 2 historical tracking across Bronze, Silver, and Gold layers.
- Built configurable ingestion and transformation frameworks using Azure SQL control tables to manage metadata, mappings, watermarks, audit logging, reconciliation, batch monitoring, and SLA tracking, enabling rapid onboarding of new countries through configuration with minimal code changes.
- Optimized large-scale Spark pipelines processing over 50 million pharmaceutical market and sales records using partitioning, broadcast joins, caching, Adaptive Query Execution (AQE), Delta OPTIMIZE, Z-ORDER, and parallel country processing to improve scalability, query performance, and production reliability.
- Reduced end-to-end reporting latency by up to 67%, accelerated new country onboarding by over 40% through metadata-driven configurations, and improved overall pipeline performance by 30–40% through Spark optimization and parallel processing while consistently meeting production SLA timelines.

Project 2: Omnichannel Retail Data Platform

Client: Leading US Retail Client | Solutionec

Tech Stack: Azure Data Factory (ADF), Azure Databricks, Azure Synapse Analytics, PySpark, Spark SQL, Delta Lake, ADLS Gen2, Azure SQL Database, Unity Catalog, Azure Event Hubs, Apache Kafka, Databricks Auto Loader, Structured Streaming, Power BI, Azure Monitor, Git

- Designed and implemented an Azure Lakehouse platform integrating batch and near real-time data from POS, e-commerce, mobile applications, marketplaces, ERP, CRM, and inventory systems into a unified retail analytics platform supporting Customer 360 and cross-channel reporting.
- Developed Azure Data Factory pipelines, Databricks PySpark notebooks, and Structured Streaming workflows for watermark-based incremental loading, Auto Loader ingestion, Delta MERGE/UPSERT operations, SCD Type 2 historical tracking, and Medallion Architecture (Bronze, Silver, Gold) data processing.
- Built scalable business transformation frameworks for customer identity resolution, sales reconciliation, inventory standardization, business rule implementation, data quality validation, and Gold layer dimensional models using Azure Databricks, Azure Synapse Analytics, and Power BI.
- Optimized analytical workloads using Azure Synapse SQL performance tuning, partition elimination, efficient indexing strategies, optimized join execution, and query tuning techniques to improve reporting performance, dashboard responsiveness, and enterprise-scale analytics.
- Implemented centralized audit, reconciliation, monitoring, and SLA tracking frameworks using Azure SQL control tables, enabling reliable batch and streaming data processing, automated operational alerts, and consistent enterprise reporting.

Project 3: Healthcare Master Data Management (MDM)

Client: Multiple Leading US Pharmaceutical Clients | Axtria

Tech Stack: Azure Data Factory (ADF), Azure Databricks, PySpark, Python, Spark SQL, Azure SQL Database, ADLS Gen2, Amazon S3, REST APIs, Linux, Azure DevOps, JIRA

- Worked as part of the Data Engineering team delivering Master Data Management (MDM) solutions for multiple US pharmaceutical clients, supporting Healthcare and Life Sciences commercial data domains including HCP, HCO, Product, Address, and Customer master data across Azure and AWS cloud platforms.
- Developed PySpark notebooks and Python scripts implementing business rules, data cleansing, standardization, referential integrity checks, email validations, null validations, duplicate detection, data profiling, and curated L3 staging datasets supporting enterprise MDM ingestion pipelines.
- Designed and optimized advanced SQL queries, database views, joins, window functions, reconciliation frameworks, and source-to-target validation processes across Landing, Pre-Staging, Change History, and Staging layers to ensure data quality, consistency, and reporting accuracy for large-scale healthcare datasets.
- Built and supported Azure Data Factory pipelines integrating Amazon S3, ADLS Gen2, REST APIs, enterprise databases, and downstream MDM platforms while automating data ingestion workflows, monitoring pipeline executions, audit frameworks,

production deployments, and operational support using Azure DevOps, JIRA, and Linux commands through PuTTY for AWS-based operational activities.

- Collaborated with enterprise MDM teams supporting inbound and outbound MDM integrations, survivorship rules, match and merge processes, crosswalk management, reference data standardization, and REST API integrations to deliver trusted master data for downstream healthcare business applications.

EDUCATION

BE in Computer Science and Engineering

Nitte Meenakshi Institute of Technology, Bengaluru

Graduated: May 2021